



Assignment 3 (13.05.2022)

Handin until: Friday, 20.05.2022, 09:00

1. [15 Points] Decomposition and MonetDB

`decompositon.sql` contains the following table:

t			
a	b	c	d
6	'a'	3	'f'
7	'b'	4	'g'
8	'c'	3	'f'
9	'd'	4	'g'
0	'e'	3	'f'

Load the table into MonetDB using:

```
mclient -l sql <dbname> <path/to/decomposition.sql>
```

- Given table $t(a,b,c,d)$, find any non-trivial non-key *Functional Dependencies (FDs)*¹.
- Based on the FDs in (a), perform a BCNF split to create BCNF tables t_1 and t_2 from t .
- Split table t_1 as well as t_2 into multiple *two-column* (binary) SQL tables whose head column carries an *id* and whose tail column carries the values of column a (or b, c, d).
Hint: At this stage, you should end up with a total of five such two-column tables.
- Formulate a MonetDB SQL query that shows that you can faithfully reconstruct t using all binary tables created in (c).
- Based on the query in (d), explain briefly why MonetDB decides to *not* pursue normalization into BCNF form (but instead directly translates the original table t into binary tables).

¹A refresher on FDs and BCNF can be found here: <https://github.com/DBatUTuebingen/db1-ws2122/blob/master/slides/db1-10.pdf>

2. [15 Points] Page Layout in PostgreSQL

Use file `documents.sql` to create table `documents(title CHAR(4), doc TEXT)`. Rows in `documents` are — due to type `TEXT` of column `doc` — of variable length and may therefore grow and exceed the free page space on `UPDATE`. How does *PostgreSQL* handle that?

Use the *PostgreSQL* extension `pageinspect`² to observe *PostgreSQL*'s behavior. Hand in all queries you used and describe your findings briefly. Proceed as follows:

(a) Use function

```
heap_page_item_attrs(get_raw_page('documents', <page>), 'documents')
```

to inspect the organization of rows on all pages of table `documents`. Next to each row's header information (`lp`: `slotno`, `lp_off`: row pointer, `lp_len`: row size, `t_ctid`: row version-chain pointer³), the function extracts the raw data of all attributes in an array `t_attrs` of type `bytea[]`.

Use function `convert(str BYTEA)→TEXT` provided in `documents.sql` to convert this attribute's data to `TEXT`.

- (b) Determine the *RID* (`page`, `slotno`) and row size of the row containing 'doc1' as well as the free space left on its page.
- (c) Perform an `UPDATE` on 'doc1' doubling the size of its `doc` column, therefore exceeding the free page space. Does the *RID* still point to the same row? How are the pages and the physical location of the row data reorganized?
- (d) Perform an `UPDATE` on 'doc2' growing its row size to more than 8 kB. The new row cannot fit into any page — even an empty one. How does *PostgreSQL* cope with that?
 - i. How does the row size of the new row compare to the size of the inserted `doc`-value?
 - ii. Read “Out-of-line, on-disk TOAST storage”⁴ of the *PostgreSQL* documentation. Explain in your own words, how “sliced bread”⁵ relates to *PostgreSQL* in terms of our current problem.
 - iii. Search the system catalog table `pg_class`⁶ to find the `relname` of the TOAST table associated with table `documents`.
 - iv. Query table `pg_toast.<TOAST_relname>` to determine, how many *chunks* (rows of the toast table) have been created to store your new `doc`-value.

²<https://www.postgresql.org/docs/current/pageinspect.html>

³See Slide 08 of Chapter 5 “Row Updates”

⁴<https://www.postgresql.org/docs/current/storage-toast.html#STORAGE-TOAST-ONDISK>

⁵<https://en.wikipedia.org/wiki/Toast>

⁶<https://www.postgresql.org/docs/current/catalog-pg-class.html>