



## Assignment 7

Handin until: Friday, 01.07.2022, 09:00

### 1. [8 Points] B<sup>+</sup>Tree - Insert




Figure 1: A B<sup>+</sup>Tree

Assume that all key values inserted into the B<sup>+</sup>Tree are unique and that the **Insert**-operation of the index does **not implement redistribution**.

Answer the following questions, each based on the *unmodified* B<sup>+</sup>Tree of Figure 1:

- Identify *four* leaf node entries ( $a, \dots, d$ ) which, when inserted in sequence (**Insert**  $a, \dots, \text{Insert } d$ ), completely fill the leaf pages.
- What is the minimum number of **Insert** operations required to increase the size of the tree by two levels?
- Take all entries of the leaf level sequence set (2,3,5,...) and **Insert** them, one by one, into a new B<sup>+</sup>Tree of order  $o = 1$ . Submit sketches of the B<sup>+</sup>Tree instances after the first three **Insert** steps, along with the final resulting B<sup>+</sup>Tree.

**Note:** You can submit your solutions as plain text, PDF or PNG image file. However, handwritten solutions are **not** allowed and will not be graded.

## 2. [12 Points] B+Tree - Maintenance

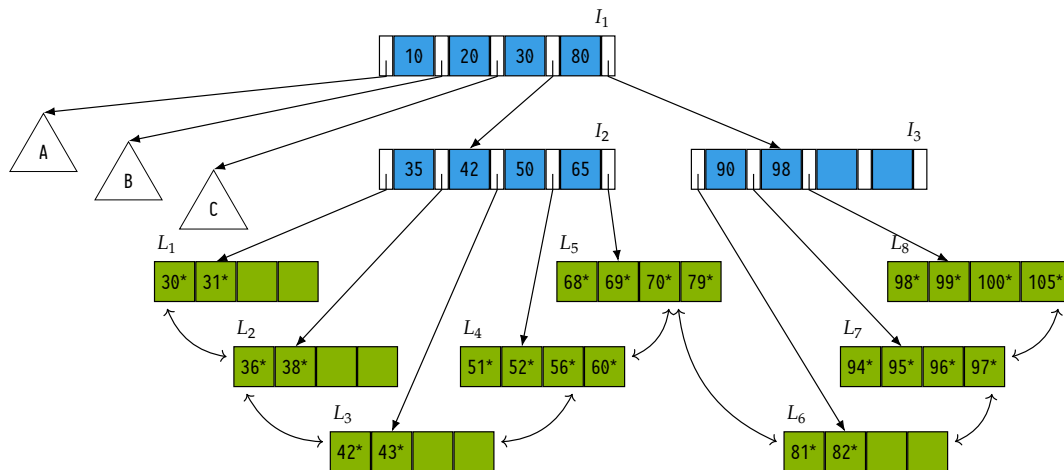


Figure 2: Another B+Tree

Assume that all key values inserted into the B+Tree are unique. However, assume both index operations, **Insert** and **Delete**, to implement **redistribution at leaf level**. Note that redistribution occurs only between direct siblings, that is, between directly connected leaf nodes that share the same parent node.

Answer the following questions, each based on the *unmodified* B+Tree of Figure 2:

- Note all nodes ( $I_j$  or  $L_k$ ) that need to be read by the RDBMS to answer the following queries:
  - “Find all records with a key value greater than 38”
  - “Delete the record with key value 81”
- Add a record with the key value 104 to the tree. Sketch the resulting tree.
- Name any key values that directly result in an increase in tree depth when inserted.
- Delete the entry with key value 36 from the tree. Sketch the resulting tree.
- The subtrees A, B and C have not been fully specified. Describe all characteristics that can still be inferred about these subtrees.

**Note:** You may hand in all tree sketches as plain text, PDF or PNG image file. Again, handwritten solutions are **not** allowed and will not be graded.

## 3. [10 Points] B+Tree - PostgreSQL

The file `btree.sql` creates and populates table

indexed (a **INT PRIMARY KEY**, b **TEXT**, c **NUMERIC(3,2)**).

Load the file with PostgreSQL and answer the following questions. Hand in all SQL queries and intermediate results used to find the final answer.

**Note:** The tasks require you to use functions `bt_metap(relname TEXT)`, `bt_page_stats(relname TEXT, blkno INT)` and `bt_page_items(relname TEXT, blkno INT)` previously mentioned in the lectures. For more information about these functions, read the documentation at

<https://www.postgresql.org/docs/current/pageinspect.html#id-1.11.7.31.6>.

- How many pages were created for the index that PostgreSQL automatically created based on the primary key?
- Write a query to find the *root node* of the B+Tree. What is its page number and what is its fan-out?
- What is the *average* fan-out of all non-leaf nodes?

- (d) Use the functions `bt_page_stats(...)` and `bt_page_items(...)` to manually traverse the B+Tree from the root to the index leaf page containing the key entry for `a = 150 000`. What are the **minimum** and **maximum key values** found on this leaf node? To which pages of relation `indexed` do the entries of that leaf node point?

**Note:** Function `bt_page_items(...)` returns a column `data` which represents the key value of an index entry. The `data` value is encoded as a hexadecimal string. For example, the value `'77 97 01 00 00 00 00 00'` is read as the hexadecimal number `0x19777` which can be converted to the decimal number `104311`.

For convenience, we have provided function `data_to_numeric(TEXT)` in `btree.sql`, to convert these `data` values to decimal numbers. It can be used as follows:

```
1 | SELECT data_to_numeric('77 97 01 00 00 00 00 00');
```

```
data_to_numeric
-----
              104311
(1 row)
```