



## Assignment 10

Handin until: Friday, 22.07.2022, 09:00

### 1. [22 Points] Join Operators in PostgreSQL

**Important Note:** Before you start working on this assignment, **disable parallelism**:

```
1 set max_parallel_workers_per_gather = 0;
2 set max_parallel_workers = 0;
```

**Nested Loop Join:** Disable Hash- and Merge Join to answer the following questions!

```
1 set enable_hashjoin = off;
2 set enable_mergejoin = off;
```

- (a) Create tables **one** and **many** as provided in **one-many.sql**. As long as no index is defined for any table, a query to join the tables will use a simple *Nested Loop Join* and will not terminate in an reasonable amount of time. Let's disable materialization first with

```
1 set enable_material = off;
```

and use **EXPLAIN** (without **ANALYZE**) to show the most naive plan for the following query **Q**:

```
1 SELECT *
2 FROM   one AS o, many AS m
3 WHERE  o.a = m.a
```

Based on the estimated rows, how often would the *Join Filter* **o.a = m.a** be evaluated?

- (b) Re-enable *materialization* and compare the new plan for query **Q** with (a).

```
1 set enable_material = on;
```

Explain why the loop order has changed and why **Materialize** is used on the **Seq Scan** of **one**, instead of **many**.

- (c) A **PRIMARY KEY** index on **one(a)** supports the *Nested Loop Join* on query **Q**. How?

```
1 ALTER TABLE one ADD CONSTRAINT one_a PRIMARY KEY (a);
2 ANALYZE;
```

Show the plan with **EXPLAIN ANALYZE** and explain it briefly.

- (d) An additional **PRIMARY KEY** index on **many(a,c)** further improves the query performance.

```
1 ALTER TABLE many ADD CONSTRAINT many_a_c PRIMARY KEY (a,c);
2 ANALYZE;
```

- Why is only one of the two indexes used, while the other table is accessed using a **Seq Scan**?
- Why is table **many** with index **many\_a\_c** (and not table **one** with **one\_a**) preferred here as an inner join table?

- (e) How does the following change to the query **Q** benefit from both indexes instead?

```
1 SELECT *
2 FROM   one AS o, many AS m
3 WHERE  o.a = m.a
4 ORDER BY m.a
```

**Hash Join:** Re-enable Hash- and Merge Join to answer the following questions!

```
1 | set enable_hashjoin = on;  
2 | set enable_mergejoin = on;
```

- (f) If available, a *Hash Join* is used to answer the equi-join query Q. Show the plan with EXPLAIN (VERBOSE, ANALYZE, BUFFERS).
- Why is table **one** (and not table **many**) chosen as the inner *build table*?
  - Why are the indexes **one\_a** and **many\_a\_c** not used to access the base tables?
- (g) *Hash Join* builds up a temporary hash table and therefore suffers when **work\_mem** is reduced. Issue `set work_mem='64kB'` (instead of default `'4MB'`) and re-execute query Q. Use the output of EXPLAIN (VERBOSE, ANALYZE, BUFFERS) to compare it with (f). Why does the performance of the *Hash Join* decrease significantly?
- (h) Since even with low **work\_mem**, some of the hash table is stored in-memory, the actual performance of (g) may depend heavily on the data distribution of the table being searched. Table **many\_skewed** in **one-many.sql** provides a variant of table **many** with a highly skewed distribution of column **a**.

Examine columns **n\_distinct**, **most\_common\_vals** and **most\_common\_freqs** in table **pg\_stats**<sup>1</sup> to show statistics about the distribution of values for both, attribute **a** in **many** and **a** in **many\_skewed**.

- Give a short comparison.

Execute the following query on **work\_mem='64kB'** and compare its plan to (g).

**Note:** You may have to execute each query twice to avoid inconsistencies in caching.

```
1 | SELECT *  
2 | FROM   one AS o, many_skewed AS m  
3 | WHERE  o.a = m.a
```

The I/O on temporary tables is reduced. To which number and why?

**Merge Join:** Disable Hash Join to answer the following question!

```
1 | set enable_hashjoin = off;
```

- (i) When we enforce *Merge Join* in (g) (Q with low **work\_mem** of `'64kB'`), it performs better than the original *Hash Join*.
- How does the *Merge Join* make use of indexes **one\_a** and **many\_a\_c**?
  - Why is memory no crucial factor here, so that *Merge Join* can outperform the *Hash Join* on **work\_mem='64kB'**?

## 2. [8 Points] External Merge Sort

Given an input table of unsorted values distributed over 16 pages. Each page covers up to two elements:

pages 1–8:	38,58	23,31	36,59	21,35	19,53	13,25	56,22	54,21
pages 9–16:	34,58	19,14	48,46	47,32	60,40	58,60	55,24	50,54

Sort the table using **External Merge Sort**, with an available working memory of **B = 3 pages**. For each *Pass* of the algorithm, write down the *output runs* written to secondary memory.

**Note:** *Pass 0* does **not** use *Replacement Sort* and returns sorted runs of size **B** pages.

<sup>1</sup><https://www.postgresql.org/docs/current/static/view-pg-stats.html>