

DB 2

07 – Expression Evaluation

Summer 2022

Torsten Grust
Universität Tübingen, Germany

1 | Q_6 — Expression Evaluation

For a large class of queries, the **CPU effort to evaluate (complex) expressions** may easily match the time spent for I/O and data access:

```
SELECT t.a * 3 - t.a * 2    AS a,  
       t.a - power(10, t.c) AS diff,  
       ceil(t.c / log(2))  AS bits  
FROM   ternary AS t;
```

Iterate over rows t , access required fields (here: $t.a$, $t.c$), evaluate (multiple) expressions per row, construct resulting row.



Using **EXPLAIN** on Q_6

EXPLAIN VERBOSE

```
SELECT t.a * 3 - t.a * 2    AS a,
       t.a - power(10, t.c) AS diff,
       ceil(t.c / log(2))  AS bits
FROM   ternary AS t;
```

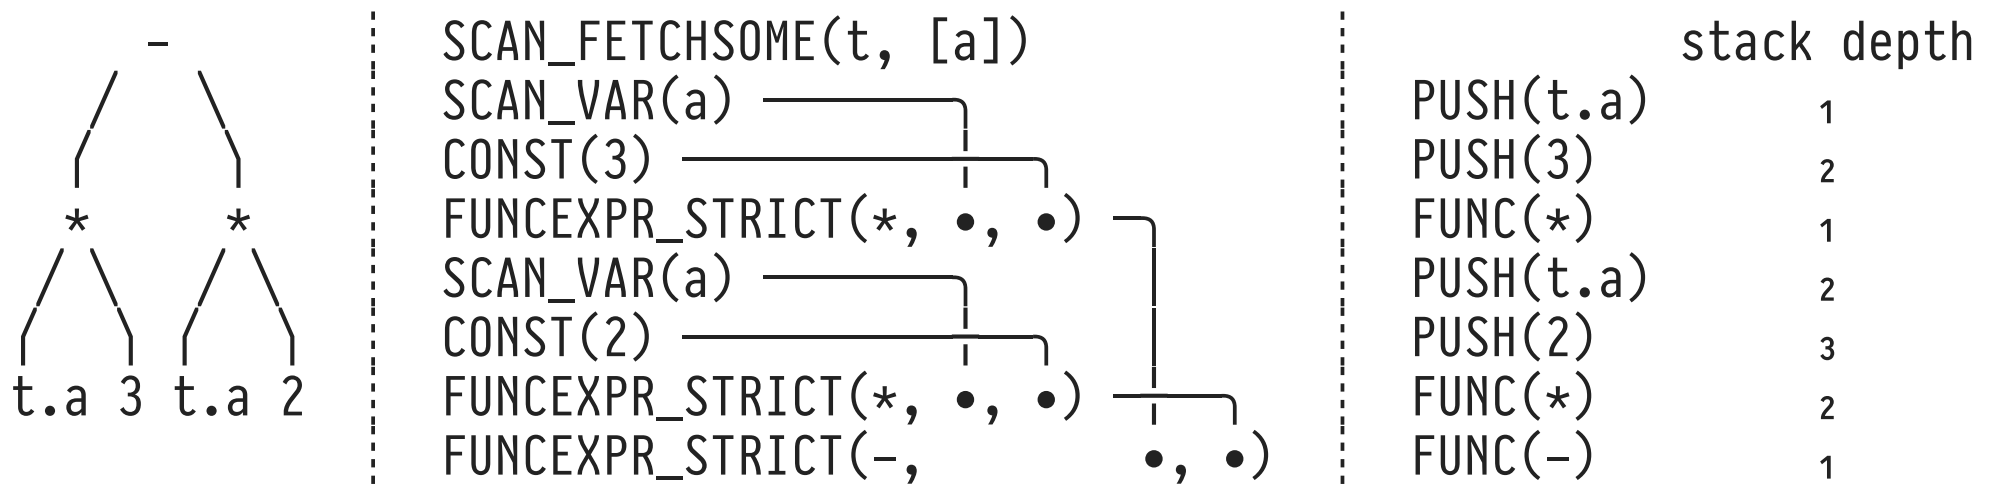
QUERY PLAN

Seq Scan on public.ternary t (cost=0.00..40.00 rows=1000 width=20)	
Output: ((a * 3) - (a * 2)),	←
((a)::double precision - power('10'::double precision, c)),	←
ceil((c / '0.301029995663981'::double precision))	←

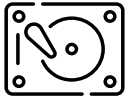
- Expressions have been parenthesized, simplified, and annotated with type casts as required by SQL semantics.

Internal Representations of $t.a * 3 - t.a * 2$

- DBMSs—just like interpreters and compilers—**transform expressions into internal representations** that facilitate simplification and evaluation:

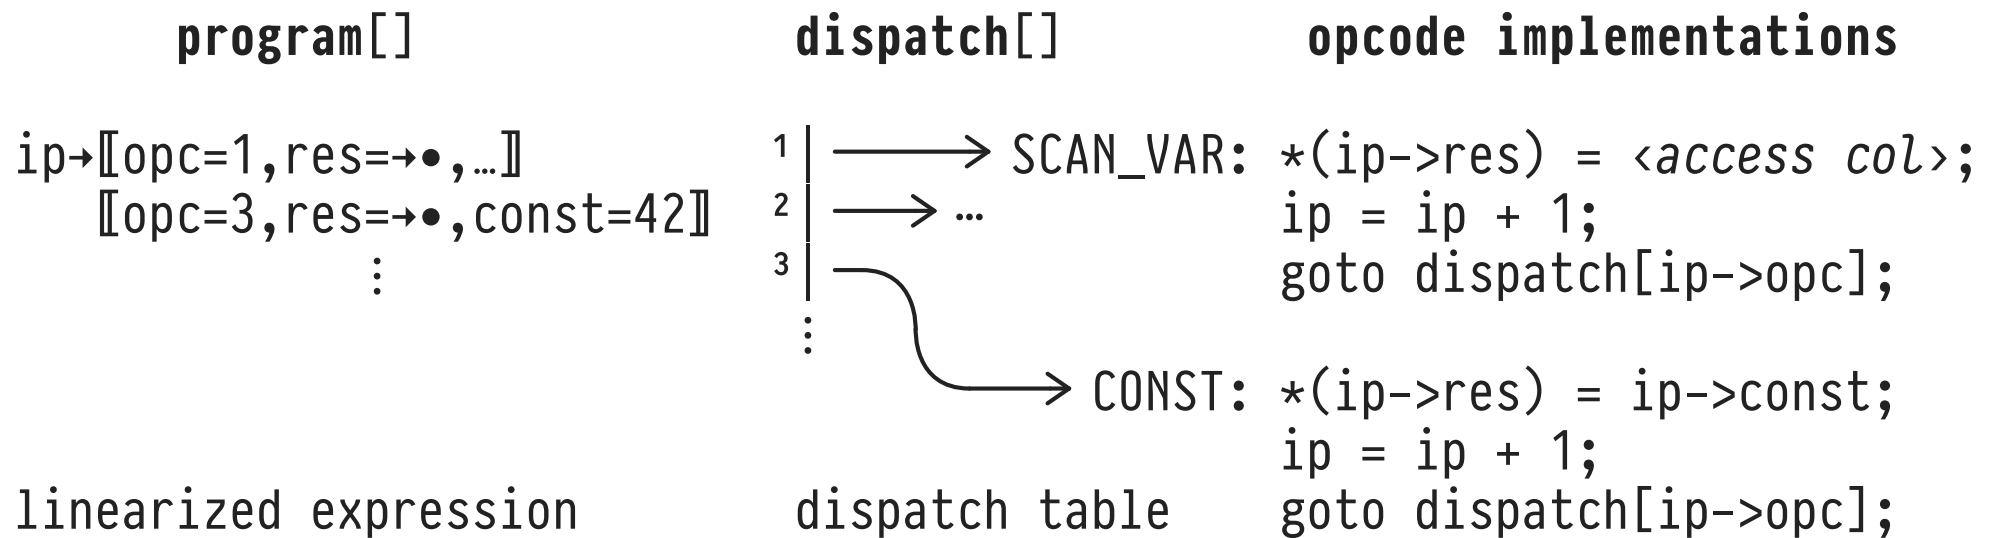


- Postorder traversal of expression tree to obtain a linearized “program”. Arg slots (\bullet) or stack push/pop.



Threaded Interpretation

PostgreSQL implements a **threaded interpreter** over linearized expressions (middle column of previous slide):



- Note: **ip**: instruction pointer, **opc**: operation code.
- Relies on support for *computed goto* (e.g., common in C).



Expression Interpretation Overhead

Overhead of expression interpretation has been found to be **massive** in DBMS (cf. the threaded interpretation vs. machine code for `t.a * 2`).

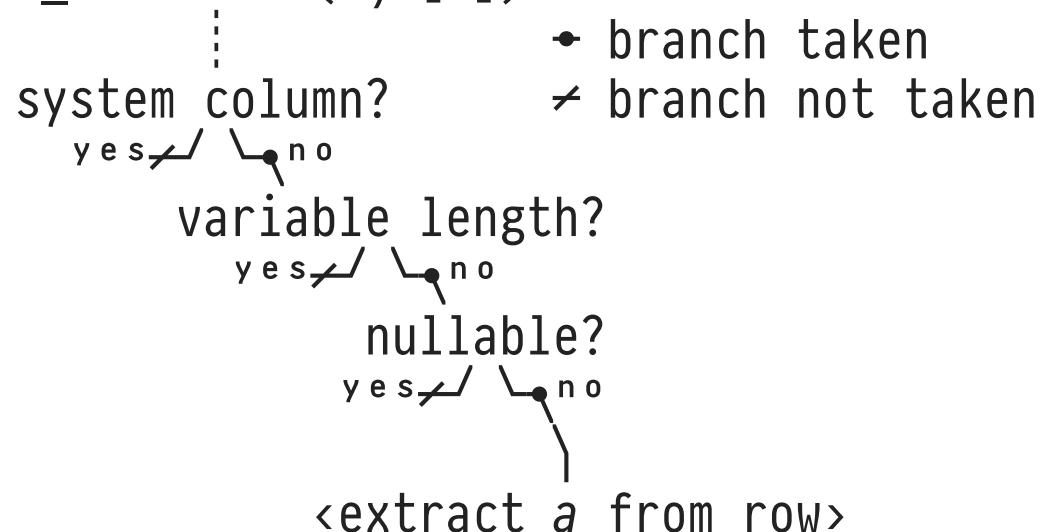
- Field access and interpretation in *hot query code path*, rediscovers same row structure and follows same opcode pointers for every row processed. Wasteful.
- 💡 Invest in **just-in-time (JIT) compilation** of expression program into machine code once, benefit for all subsequent rows.
 - **NB.** LLVM-based support for JIT compilation of expressions has been added to PostgreSQL since v11.



JIT: Turn Run-time into Compile-time Decisions

- PostgreSQL's interpreter is very generic, prepared to handle corner cases, exceptions, and extensions. Leads to branch-heavy routines in the interpreter's hot code path.
- Expression compilation creates **query-specific code**:

SCAN_FETCHSOME(t , $[a]$)



- Access field a of a row.
- Interpreter follows **same code path** for all rows, possibly millions of times.
- Generate specific “ a access code” along the ⚡ path.



Expression Compilation: When (Not) to JIT?

- JITing involves code generation, optimization, IR emission (LLVM bitcode), and translation to native code.
- JIT effort adds to the query planning time. Be careful not to penalize queries Q that are cheap to begin with:

Assume runtime of query Q reduced by 20% due to JIT compilation:



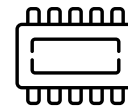
2 : Q_6 — Column-Based Expression Evaluation



```
SELECT t.a * 3 - t.a * 2    AS a,  
       t.a - power(10, t.c) AS diff,  
       ceil(t.c / log(2))  AS bits  
FROM   ternary AS t;
```

MonetDB compiles expressions into sequences of MAL operations. Like data processing, expression evaluation is column-oriented (as opposed to row-by-row).

- We will find that this vector-based evaluation mode fits modern CPU architecture particularly well.

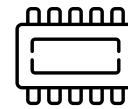


Using EXPLAIN on Q_6

```
sql> EXPLAIN SELECT t.a * 3 - t.a * 2 AS a,
                ceil(t.c / log10(2)) AS bits
                FROM ternary AS t;

:
ternary :bat[:oid] := sql.tid(sql, "sys", "ternary");
c0       :bat[:dbl] := sql.bind(sql, "sys", "ternary", "c", 0:int);
c        := algebra.projection(ternary, c0);
e1       :bat[:dbl] := batcalc./(c, 0.3010299956639812:dbl);
e2       :bat[:dbl] := batmath.ceil(e1);                      ← result column 'bits'
a0       :bat[:int] := sql.bind(sql, "sys", "ternary", "a", 0:int);
a        := algebra.projection(ternary, a0);
e3       :bat[:lng] := batcalc.lng(a);                        ← cast to type lng
e4       :bat[:lng] := batcalc.*(e3, 3:bte);
e5       :bat[:lng] := batcalc.*(e3, 2:bte);
e6       :bat[:lng] := batcalc.-(e4, e5);                      ← result column 'a'
:
```

- MAL ops `batcalc.⊗` accept two BATs or one BAT + one scalar (like `2:bte`, `3:bte`, `0.301...:dbl` \equiv `log10(2)`).

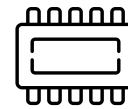


Column-Based “Zip” Semantics

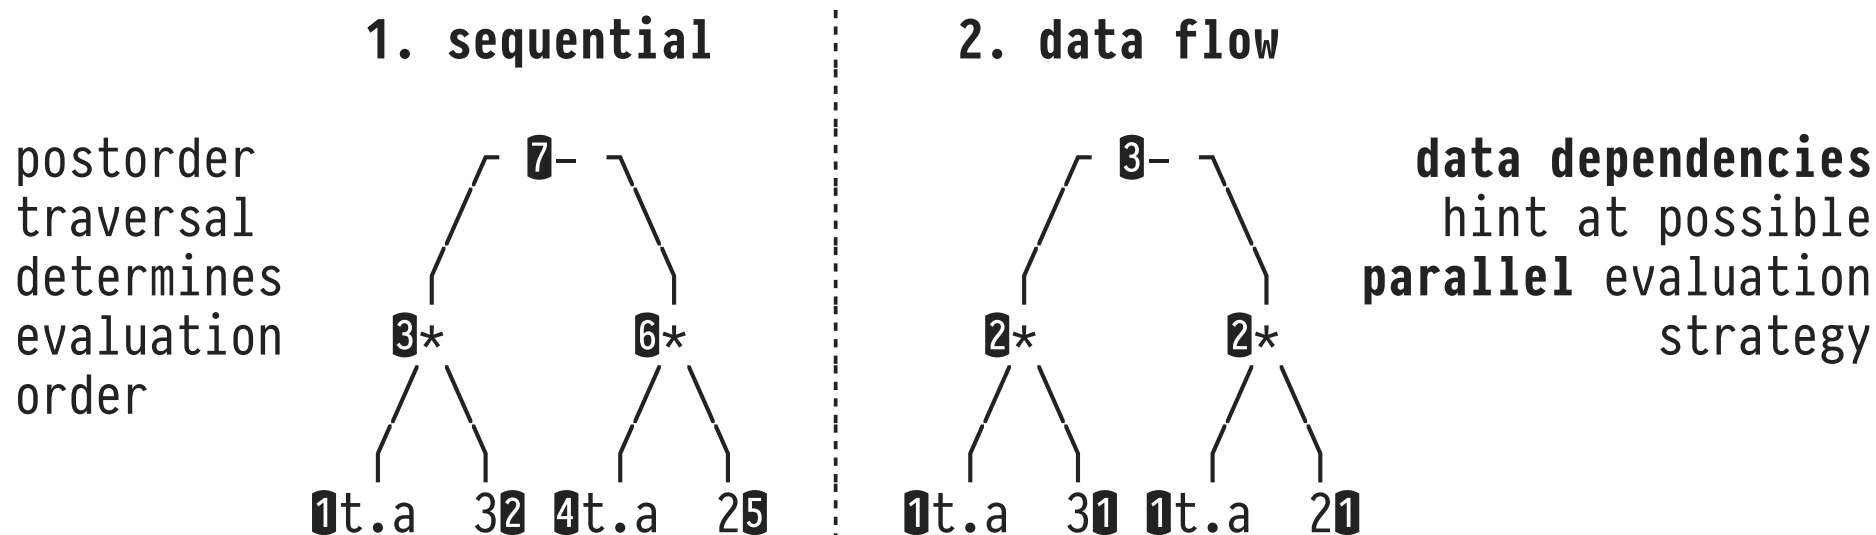
Operators `batcalc.⊗` merge the tails of two synchronized BATs using binary operator \otimes , yields a new BAT:

x			batcalc.-(x,y)			y	
head	tail		head	tail		head	tail
0@0	x_0	→	0@0	$x_0 - y_0$	←	0@0	y_0
1@0	x_1	→	1@0	$x_1 - y_1$	←	1@0	y_1
2@0	nil	→	2@0	nil	←	2@0	y_2
3@0	x_3	→	3@0	$x_3 - y_3$	←	3@0	y_3

- `batcalc.⊗` contains checks for arithmetic exceptions (overflow, divide by 0). Also: `nil ⊗ x = x ⊗ nil = nil`.




MAL: Sequential Execution vs. Data Flow



1. Order of assignment to temporary result BATs e_i follows postorder traversal of expression tree.
2. Spawn CPU threads to evaluate data-independent subexpressions in // (see MonetDB's [dataflow](#) optimizer).

batcalc.⊗: Column-Based Operator Implementations (1)




MonetDB supplies type- and \otimes -specific implementations of MAL operations (code generation via C preprocessor macros):

```
/* batcalc.-(left:bat[:lng], right:bat[:lng]):bat[:lng] */
                                     ↑
int i, j, k;
int nils = 0;

for (i = start, j = start*1, k = start; k < end; i += 1, j += 1, k += 1) {
    /* nil checking */
    if (is_lng_nil(left[i]) || is_lng_nil(right[j])) {
        result[k] = lng_nil;
        nils++;
    } else {
        /* omitted: overflow checking (abort on error or emit nil) */
        result[k] = left[i] - right[j];
    }
}
```

batcalc.⊗: Column-Based Operator Implementations (2)



MonetDB supplies type- and \otimes -specific implementations of MAL operations (code generation via C preprocessor macros):

```
/* batcalc.-(left:bat[:lng], right:lng):bat[:lng] */
                                     ↑
int i, j, k;
int nils = 0;

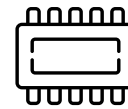
for (i = start, j = start*0, k = start; k < end; i += 1, j += 0, k += 1) {
    /* nil checking */
    if (is_lng_nil(left[i]) || is_lng_nil(right[j])) {
        result[k] = lng_nil;
        nils++;
    } else {
        /* omitted: overflow checking (abort on error or emit nil) */
        result[k] = left[i] - right[j];
    }
}
```

3 : Column-Based Operators vs. Expression Interpretation

Expression evaluation through column-based operator and row-wise interpretation compared:

Column-Based (MonetDB)	Row-Wise (PostgreSQL)
zero degrees of freedom instruction locality optimizable tight loops <ul style="list-style-type: none"> • loop pipelining • blocking • loop unrolling data parallelism full materialization	variable-width rows w/ fields of various types computed goto, long code paths complex control flow, code in many functions <ul style="list-style-type: none"> • unpredictable branches focus on single row row-by-row result generation

- Compilers **optimize tight code loops** inside MAL operators.
- CPUs offer wide registers and instructions to exploit **data //ism** (SIMD: *single instruction, multiple data*).



Compiling Tight Loops (cf. MAL Operators)

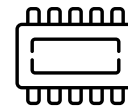
Inspect Intel® x86 code generated by LLVM's C compiler `clang` for MonetDB's routine `BATcalcsb` (`batcalc.-`), simplified:

```
#define SIZE 1024

void BATcalcsb(int *left, int *right, int *result)
{
    int i, j, k;

    for (i = j = k = 0; k < SIZE; i += 1, j += 1, k += 1) {
        result[k] = left[i] - right[j];
    }
}
```

- Arrays `left`, `right`/`result` represent input/output BATs.



Assembly Code for Simple Tight Loop


Uses `clang` (options `-O2 -fno-vectorize -fno-unroll-loops`).

- Register assignment:

`left: %rdi, right: %rsi, result: %rdx, i/j/k: %rax`

```
BATcalcsb:  
    xorl %eax, %eax                # %rax ←32 0  
loop:  
    movl (%rdi,%rax,4), %ecx        # %ecx ←32 mem[%rdi + %rax]  
    subl (%rsi,%rax,4), %ecx        # %ecx ←32 %ecx -32 mem[%rsi + %rax]  
    movl %ecx, (%rdx,%rax,4)        # mem[%rdx + %rax] ←32 %ecx  
    addq $1, %rax  
    cmpq $1024, %rax               # 1024 loop iterations  
    jne  loop                      # exit if %rax = 1024  
    retq
```

- **NB.** One loop exit test per array element computed.



(Explicit) Loop Unrolling

- Manually perform **loop unrolling** to
 1. improve the ratio (*useful work*) / (*loop exit test*),
 2. expose independent work that may be executed in `//`:


```
void BATcalcsb(int *left, int *right, int *result)
{
    int i, j, k;

    for (i = j = k = 0; k < SIZE; i += 4, j += 4, k += 4) {
        result[k] = left[i] - right[j];
        result[k+1] = left[i+1] - right[j+1];
        result[k+2] = left[i+2] - right[j+2];
        result[k+3] = left[i+3] - right[j+3];
    }
}
```

↓ ↓ ↓

} independent, execute in
any order or even in //

- **NB.** Needs code to handle the case `SIZE mod 4 ≠ 0`.



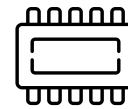
Loop Unrolling

Compiler `clang` (options `-O2 -fno-vectorize -funroll-loops`) unrolls four loop iterations (easy for CPU to //ize):

```

BATcalcsb:
    xorl %eax, %eax                # i/j/k ←32 0
loop:
    movl (%rdi,%rax,4), %ecx        # %ecx ←32 left[i]
    subl (%rsi,%rax,4), %ecx        # %ecx ←32 %ecx -32 right[j]
    movl %ecx, (%rdx,%rax,4)        # result[k] ←32 %ecx
    movl 4(%rdi,%rax,4), %ecx       # %ecx ←32 left[i+1]
    subl 4(%rsi,%rax,4), %ecx       # %ecx ←32 %ecx -32 right[j+1]
    movl %ecx, 4(%rdx,%rax,4)       # result[k+1] ←32 %ecx
    movl 8(%rdi,%rax,4), %ecx       # :
    subl 8(%rsi,%rax,4), %ecx
    movl %ecx, 8(%rdx,%rax,4)
    movl 12(%rdi,%rax,4), %ecx
    subl 12(%rsi,%rax,4), %ecx
    movl %ecx, 12(%rdx,%rax,4)
    addq $4, %rax                  # }
    cmpq $1024, %rax              # } 1024 / 4 = 256 loop iterations
    jne loop                      # exit if %rax = 1024
    retq

```



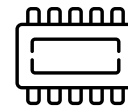
Data-Parallelism Through SIMD

result[] left[] right[]

0	[]	← 1 2 8	[]	- 1 2 8	[]
1	[]		[]		[]
2	[]		[]		[]
3	[]		[]		[]
4	[]	← 1 2 8	[]	- 1 2 8	[]
5	[]		[]		[]
6	[]		[]		[]
7	[]		[]		[]
8	[]		[]		[]

- Read/compute/write four array elements (of width 4×32 bits = 128 bits) at a time in **data-parallel** fashion.
- Relies on SIMD register and instructions (e.g., Intel® SSE registers `%xmmi` and instruction `move double quad word`)

- ⚠ Requires care if
 - arrays `result[]` and `left[]/right[]` overlap in memory,
 - residual array elements (see []) are to be processed.

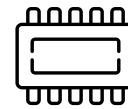


Data-Parallelism Through SIMD (Array Overlap)

Compiler `clang` (options `-O2 -fvectorize`) uses SIMD registers and instructions.

- Extra prelude code checks for **array overlap**. If so, jumps to non-vectorized (yet unrolled) version of code.
- Declare function arguments via `restrict` to inform C compiler that arrays won't overlap:

```
void BATcalcsb(int *restrict left, ..., int *restrict result)
{
    int i, j, k;
    ⋮
}
```



Data-Parallelism Through SIMD (Main Loop)

Process 16 elements per iteration (SIMD + 2 loops unrolled, assumes no overlap of arrays `result[]` and `left[]/right[]`):

```

xorl %eax, %eax                                4 × 32 bits = 128 bits wide
loop:
movdqu (%rdi,%rax,4), %xmm0                    # %xmm0 ←128 left[i+0...i+3]
movdqu 16(%rdi,%rax,4), %xmm1                  # %xmm1 ←128 left[i+4...i+7]
movdqu (%rsi,%rax,4), %xmm2                    # %xmm2 ←128 right[i+0...i+3]
psubd %xmm2, %xmm0                             # %xmm0 ←128 %xmm0 -128 %xmm2
movdqu 16(%rsi,%rax,4), %xmm2                  # %xmm2 ←128 right[i+4...i+7]
psubd %xmm2, %xmm1                             # %xmm1 ←128 %xmm1 -128 %xmm2
movdqu %xmm0, (%rdx,%rax,4)                     # result[i+0...i+3] ←128 %xmm0
movdqu %xmm1, 16(%rdx,%rax,4)                   # result[i+4...i+7] ←128 %xmm1
movdqu 32(%rdi,%rax,4), %xmm0                    # %xmm0 ←128 left[i+8...i+11]
movdqu 48(%rdi,%rax,4), %xmm1                    # %xmm1 ←128 left[i+12...i+15]
movdqu 32(%rsi,%rax,4), %xmm2                    # %xmm2 ←128 right[i+8...i+11]
psubd %xmm2, %xmm0                             # %xmm0 ←128 %xmm0 -128 %xmm2
movdqu 48(%rsi,%rax,4), %xmm2                    # %xmm2 ←128 right[i+12...i+15]
psubd %xmm2, %xmm1                             # %xmm1 ←128 %xmm1 -128 %xmm2
movdqu %xmm0, 32(%rdx,%rax,4)                   # result[i+8...i+11] ←128 %xmm0
movdqu %xmm1, 48(%rdx,%rax,4)                   # result[i+12...i+15] ←128 %xmm1
addq $16, %rax                                  # }
cmpq $1024, %rax                               # } 1024 / 16 = 64 iterations
jne loop                                       # exit if %rax = 1024
:                                              # (non-vectorized code not shown)

```

..... loop #n

..... loop #n+1